

IMDB

July 3, 2024

<div class="alert alert-success"; style="border-left: 7px solid green"> Reviewer's comment, v. 2

Great! You've done a great job and now your project has been accepted.

Thank you for your work and I wish you success in the following projects!

Hello Dhruval

My name is Dima, and I will be reviewing your project.

You will find my comments in coloured cells marked as 'Reviewer's comment'. The cell colour will vary based on the contents - I am explaining it further below.

Note: Please do not remove or change my comments - they will help me in my future reviews and will make the process smoother for both of us.

You are also very welcome to leave your comments / describe the corrections you've done / ask me questions, marking them with a different colour. You can use the example below:

<div class="alert alert-info"; style="border-left: 7px solid blue"> Student's comment

0.1 Basic Python - Project

Review summary

Dhruval, thanks for submitting the project. You've done a very good job and I enjoyed reviewing it.

- You completed all the tasks.
- Your code was optimal and easy to read.
- You wrote your own functions.

There is only one critical comment that need to be corrected. You will find it in the red-colored cell in relevant section. If you have any questions please write them when you return your project.

I'll be looking forward to getting your updated notebook.

0.2 Introduction

In this project, you will work with data from the entertainment industry. You will study a dataset with records on movies and shows. The research will focus on the "Golden Age" of television, which began in 1999 with the release of *The Sopranos* and is still ongoing.

The aim of this project is to investigate how the number of votes a title receives impacts its ratings. The assumption is that highly-rated shows (we will focus on TV shows, ignoring movies) released during the “Golden Age” of television also have the most votes.

0.2.1 Stages

Data on movies and shows is stored in the `/datasets/movies_and_shows.csv` file. There is no information about the quality of the data, so you will need to explore it before doing the analysis.

First, you’ll evaluate the quality of the data and see whether its issues are significant. Then, during data preprocessing, you will try to account for the most critical problems.

Your project will consist of three stages: 1. Data overview 2. Data preprocessing 3. Data analysis

0.3 Stage 1. Data overview

Open and explore the data.

You’ll need `pandas`, so import it.

```
[2]: # importing pandas
import pandas as pd
```

Read the `movies_and_shows.csv` file from the `datasets` folder and save it in the `df` variable:

```
[3]: # reading the files and storing them to df
df = pd.read_csv('/datasets/movies_and_shows.csv')
```

Print the first 10 table rows:

```
[4]: # obtaining the first 10 rows from the df table
# hint: you can use head() and tail() in Jupyter Notebook without wrapping them
into print()
print(df.head(10))
```

	name	Character	role	TITLE	Type	\
0	Robert De Niro	Travis Bickle	ACTOR	Taxi Driver	MOVIE	
1	Jodie Foster	Iris Steensma	ACTOR	Taxi Driver	MOVIE	
2	Albert Brooks	Tom	ACTOR	Taxi Driver	MOVIE	
3	Harvey Keitel	Matthew 'Sport' Higgins	ACTOR	Taxi Driver	MOVIE	
4	Cybill Shepherd	Betsy	ACTOR	Taxi Driver	MOVIE	
5	Peter Boyle	Wizard	ACTOR	Taxi Driver	MOVIE	
6	Leonard Harris	Senator Charles Palantine	ACTOR	Taxi Driver	MOVIE	
7	Diahnne Abbott	Concession Girl	ACTOR	Taxi Driver	MOVIE	
8	Gino Ardito	Policeman at Rally	ACTOR	Taxi Driver	MOVIE	
9	Martin Scorsese	Passenger Watching Silhouette	ACTOR	Taxi Driver	MOVIE	

	release Year	genres	imdb score	imdb votes
0	1976	['drama', 'crime']	8.2	808582.0
1	1976	['drama', 'crime']	8.2	808582.0
2	1976	['drama', 'crime']	8.2	808582.0

3	1976	['drama', 'crime']	8.2	808582.0
4	1976	['drama', 'crime']	8.2	808582.0
5	1976	['drama', 'crime']	8.2	808582.0
6	1976	['drama', 'crime']	8.2	808582.0
7	1976	['drama', 'crime']	8.2	808582.0
8	1976	['drama', 'crime']	8.2	808582.0
9	1976	['drama', 'crime']	8.2	808582.0

Obtain the general information about the table with one command:

```
[5]: # obtaining general information about the data in df
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 85579 entries, 0 to 85578
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0    name            85579 non-null  object
1   Character       85579 non-null  object
2   r0le            85579 non-null  object
3   TITLE           85578 non-null  object
4   Type            85579 non-null  object
5   release Year    85579 non-null  int64
6   genres          85579 non-null  object
7   imdb sc0re     80970 non-null  float64
8   imdb v0tes     80853 non-null  float64
dtypes: float64(2), int64(1), object(6)
memory usage: 5.9+ MB
```

<div class="alert alert-success"; style="border-left: 7px solid green"> Reviewer's comment, v. 1
Great - you've used a comprehensive set of methods to have a first look at the data.

The table contains nine columns. The majority store the same data type: object. The only exceptions are 'release Year' (int64 type), 'imdb sc0re' (float64 type) and 'imdb v0tes' (float64 type). Scores and votes will be used in our analysis, so it's important to verify that they are present in the dataframe in the appropriate numeric format. Three columns ('TITLE', 'imdb sc0re' and 'imdb v0tes') have missing values.

According to the documentation: - 'name' — actor/director's name and last name - 'Character' — character played (for actors) - 'r0le' — the person's contribution to the title (it can be in the capacity of either actor or director) - 'TITLE' — title of the movie (show) - 'Type' — show or movie - 'release Year' — year when movie (show) was released - 'genres' — list of genres under which the movie (show) falls - 'imdb sc0re' — score on IMDb - 'imdb v0tes' — votes on IMDb

We can see three issues with the column names: 1. Some names are uppercase, while others are lowercase. 2. There are names containing whitespace. 3. A few column names have digit '0' instead of letter 'o'.

0.3.1 Conclusions

Each row in the table stores data about a movie or show. The columns can be divided into two categories: the first is about the roles held by different people who worked on the movie or show (role, name of the actor or director, and character if the row is about an actor); the second category is information about the movie or show itself (title, release year, genre, imdb figures).

It's clear that there is sufficient data to do the analysis and evaluate our assumption. However, to move forward, we need to preprocess the data.

<div class="alert alert-success"; style="border-left: 7px solid green"> Reviewer's comment, v. 1

Please note that it is highly recommended to add a conclusion / summary after each section and describe briefly your observations and / or major outcomes of the analysis.

0.4 Stage 2. Data preprocessing

Correct the formatting in the column headers and deal with the missing values. Then, check whether there are duplicates in the data.

```
[6]: # the list of column names in the df table
df.columns
```

```
[6]: Index([' name', 'Character', 'r0le', 'TITLE', ' Type', 'release Year',
         'genres', 'imdb sc0re', 'imdb v0tes'],
         dtype='object')
```

Change the column names according to the rules of good style: * If the name has several words, use snake_case * All characters must be lowercase * Remove whitespace * Replace zero with letter 'o'

```
[7]: # renaming columns
df = df.rename(columns={' name': 'name', 'Character': 'character',
                        'r0le': 'role', 'TITLE': 'title', ' Type': 'type',
                        'release Year': 'release_year',
                        'imdb sc0re': 'imdb_score', 'imdb v0tes': 'imdb_voted'
                        })
```

<div class="alert alert-success"; style="border-left: 7px solid green"> Reviewer's comment, v. 1

This is a good way to rename the columns.

Check the result. Print the names of the columns once more:

```
[8]: # checking result: the list of column names
print(df.columns)
```

```
Index(['name', 'character', 'role', 'title', 'type', 'release_year', 'genres',
      'imdb_score', 'imdb_voted'],
      dtype='object')
```

0.4.1 Missing values

First, find the number of missing values in the table. To do so, combine two **pandas** methods:

```
[9]: # calculating missing values
print(df.isna().sum())
```

```
name          0
character     0
role          0
title         1
type          0
release_year  0
genres        0
imdb_score    4609
imdb_voted    4726
dtype: int64
```

<div class="alert alert-success"; style="border-left: 7px solid green"> Reviewer's comment, v. 1

The `isna()` method is selected to find the missing values, it's great!

Not all missing values affect the research: the single missing value in 'title' is not critical. The missing values in columns 'imdb_score' and 'imdb_votes' represent around 6% of all records (4,609 and 4,726, respectively, of the total 85,579). This could potentially affect our research. To avoid this issue, we will drop rows with missing values in the 'imdb_score' and 'imdb_votes' columns.

```
[10]: # dropping rows where columns with title, scores and votes have missing values
df = df.dropna(axis = 'rows')
```

<div class="alert alert-success"; style="border-left: 7px solid green"> Reviewer's comment, v. 1

Perfect!

Make sure the table doesn't contain any more missing values. Count the missing values again.

```
[11]: # counting missing values
print(df.isna().sum())
```

```
name          0
character     0
role          0
title         0
type          0
release_year  0
genres        0
imdb_score    0
imdb_voted    0
dtype: int64
```

0.4.2 Duplicates

Find the number of duplicate rows in the table using one command:

```
[12]: # counting duplicate rows

print(df.duplicated().sum())
```

6994

Review the duplicate rows to determine if removing them would distort our dataset.

```
[13]: # Produce table with duplicates (with original rows included) and review last 5
      ↪rows
duplicates_rows = df[df.duplicated(keep = False)]
print(duplicates_rows.tail(5))
```

	name	character	role	title	type	\
85569	Jessica Cediel	Liliana Navarro	ACTOR	Lokillo	MOVIE	
85570	Javier Gardeazabal	Agustín "Peluca" Ortiz	ACTOR	Lokillo	MOVIE	
85571	Carla Giraldo	Valery Reinoso	ACTOR	Lokillo	MOVIE	
85572	Ana María Sánchez	Lourdes	ACTOR	Lokillo	MOVIE	
85577	Isabel Gaona	Cacica	ACTOR	Lokillo	MOVIE	

	release_year	genres	imdb_score	imdb_voted
85569	2021	['comedy']	3.8	68.0
85570	2021	['comedy']	3.8	68.0
85571	2021	['comedy']	3.8	68.0
85572	2021	['comedy']	3.8	68.0
85577	2021	['comedy']	3.8	68.0

There are two clear duplicates in the printed rows. We can safely remove them. Call the `pandas` method for getting rid of duplicate rows:

```
[14]: # removing duplicate rows
df = df.drop_duplicates().reset_index(drop=True)
```

Check for duplicate rows once more to make sure you have removed all of them:

```
[15]: # checking for duplicates
print(df.duplicated())
print(df.duplicated().sum())
```

```
0      False
1      False
2      False
3      False
4      False
...
73854  False
```

```
73855    False
73856    False
73857    False
73858    False
Length: 73859, dtype: bool
0
```

<div class="alert alert-success"; style="border-left: 7px solid green"> Reviewer's comment, v. 1
Great, you found and removed the duplicates. And did very thorough checks to make sure the duplicates are gone.

Now get rid of implicit duplicates in the 'type' column. For example, the string 'SHOW' can be written in different ways. These kinds of errors will also affect the result.

Print a list of unique 'type' names, sorted in alphabetical order. To do so: * Retrieve the intended dataframe column * Apply a sorting method to it * For the sorted column, call the method that will return all unique column values

```
[36]: # viewing unique type names
print(df['type'].sort_values().unique())
```

```
['MOVIE' 'movies' 'show' 'the movie']
```

<div class="alert alert-danger"; style="border-left: 7px solid red"> Reviewer's comment, v. 1

Please note, that according to the technical task it was asked:

- For the **sorted** column, call the method that will return all unique column values

So, we should use sorting here

<div class="alert alert-success"; style="border-left: 7px solid green"> Reviewer's comment, v. 2

Now it's perfect!

Look through the list to find implicit duplicates of 'show' ('movie' duplicates will be ignored since the assumption is about shows). These could be names written incorrectly or alternative names of the same genre.

You will see the following implicit duplicates: * 'shows' * 'SHOW' * 'tv show' * 'tv shows' * 'tv series' * 'tv'

To get rid of them, declare the function `replace_wrong_show()` with two parameters: * `wrong_shows_list`= — the list of duplicates * `correct_show`= — the string with the correct value

The function should correct the names in the 'type' column from the `df` table (i.e., replace each value from the `wrong_shows_list` list with the value in `correct_show`).

```
[22]: # function for replacing implicit duplicates

def replace_wrong_show(wrong_shows_list, correct_show):
    df['type'] = df['type'].replace(wrong_shows_list, correct_show)
```

Call `replace_wrong_show()` and pass it arguments so that it clears implicit duplicates and replaces them with `SHOW`:

```
[23]: # removing implicit duplicates
wrong_shows_list = ['shows', 'SHOW', 'tv show', 'tv shows', 'tv series', 'tv']
correct_show = 'show'
replace_wrong_show(wrong_shows_list, correct_show)
```

<div class="alert alert-success"; style="border-left: 7px solid green"> Reviewer's comment, v. 1
Yes, this is what was needed!

Make sure the duplicate names are removed. Print the list of unique values from the `'type'` column:

```
[24]: # viewing unique genre names
print(df['type'].unique())
```

```
['MOVIE' 'the movie' 'show' 'movies']
```

0.4.3 Conclusions

We detected three issues with the data:

- Incorrect header styles
- Missing values
- Duplicate rows and implicit duplicates

The headers have been cleaned up to make processing the table simpler.

All rows with missing values have been removed.

The absence of duplicates will make the results more precise and easier to understand.

Now we can move on to our analysis of the prepared data.

0.5 Stage 3. Data analysis

Based on the previous project stages, you can now define how the assumption will be checked. Calculate the average amount of votes for each score (this data is available in the `imdb_score` and `imdb_votes` columns), and then check how these averages relate to each other. If the averages for shows with the highest scores are bigger than those for shows with lower scores, the assumption appears to be true.

Based on this, complete the following steps:

- Filter the dataframe to only include shows released in 1999 or later.
- Group scores into buckets by rounding the values of the appropriate column (a set of 1-10 integers will help us make the outcome of our calculations more evident without damaging the quality of our research).
- Identify outliers among scores based on their number of votes, and exclude scores with few votes.
- Calculate the average votes for each score and check whether the assumption matches the results.

To filter the dataframe and only include shows released in 1999 or later, you will take two steps. First, keep only titles published in 1999 or later in our dataframe. Then, filter the table to only contain shows (movies will be removed).

```
[25]: # using conditional indexing modify df so it has only titles released after
      ↪1999 (with 1999 included)
      # give the slice of dataframe new name
      after_1999 = df[df['release_year'] >= 1999]
      print(after_1999)
```

	name	character	role	title	type \
1661	Jeff Probst	Himself - Host	ACTOR	Survivor	show
1952	Benicio del Toro	Franky Four Fingers	ACTOR	Snatch	MOVIE
1953	Dennis Farina	Cousin Avi	ACTOR	Snatch	MOVIE
1954	Vinnie Jones	Bullet Tooth Tony	ACTOR	Snatch	MOVIE
1955	Brad Pitt	Mickey O'Neil	ACTOR	Snatch	MOVIE
...
73854	A??da Morales	Maritza	ACTOR	Lokillo	the movie
73855	Adelaida Buscato	Mar??a Paz	ACTOR	Lokillo	the movie
73856	Luz Stella Luengas	Karen Bayona	ACTOR	Lokillo	the movie
73857	In??s Prieto	Fanny	ACTOR	Lokillo	the movie
73858	Julian Gaviria	unknown	DIRECTOR	Lokillo	the movie

	release_year	genres	imdb_score	imdb_voted
1661	2000	['reality']	7.4	24687.0
1952	2000	['crime', 'comedy']	8.3	841435.0
1953	2000	['crime', 'comedy']	8.3	841435.0
1954	2000	['crime', 'comedy']	8.3	841435.0
1955	2000	['crime', 'comedy']	8.3	841435.0
...
73854	2021	['comedy']	3.8	68.0
73855	2021	['comedy']	3.8	68.0
73856	2021	['comedy']	3.8	68.0
73857	2021	['comedy']	3.8	68.0
73858	2021	['comedy']	3.8	68.0

[69881 rows x 9 columns]

```
[26]: # repeat conditional indexing so df has only shows (movies are removed as
      ↪result)
      only_shows = after_1999[after_1999['type'] == 'show']
      only_shows
```

```
[26]:
```

	name	character	role	title \
1661	Jeff Probst	Himself - Host	ACTOR	Survivor
2073	Mayumi Tanaka	Monkey D. Luffy (voice)	ACTOR	One Piece
2074	Kazuya Nakai	Roronoa Zoro (voice)	ACTOR	One Piece

2075	Akemi Okamura	Nami (voice)	ACTOR	One Piece
2076	Kappei Yamaguchi	Usopp (voice)	ACTOR	One Piece
...
73831	Maneerat Kam-Uan	Ae	ACTOR	Let's Eat
73832	Rudklao Amratisha	unknown	ACTOR	Let's Eat
73833	Jaturong Mokjok	unknown	ACTOR	Let's Eat
73834	Pisamai Wilaisak	unknown	ACTOR	Let's Eat
73835	Sarawut Wichiensarn	unknown	DIRECTOR	Let's Eat

	type	release_year	genres \
1661	show	2000	['reality']
2073	show	1999	['animation', 'action', 'comedy', 'drama', 'fa...
2074	show	1999	['animation', 'action', 'comedy', 'drama', 'fa...
2075	show	1999	['animation', 'action', 'comedy', 'drama', 'fa...
2076	show	1999	['animation', 'action', 'comedy', 'drama', 'fa...
...
73831	show	2021	['drama', 'comedy']
73832	show	2021	['drama', 'comedy']
73833	show	2021	['drama', 'comedy']
73834	show	2021	['drama', 'comedy']
73835	show	2021	['drama', 'comedy']

	imdb_score	imdb_voted
1661	7.4	24687.0
2073	8.8	117129.0
2074	8.8	117129.0
2075	8.8	117129.0
2076	8.8	117129.0
...
73831	8.2	5.0
73832	8.2	5.0
73833	8.2	5.0
73834	8.2	5.0
73835	8.2	5.0

[13430 rows x 9 columns]

The scores that are to be grouped should be rounded. For instance, titles with scores like 7.8, 8.1, and 8.3 will all be placed in the same bucket with a score of 8.

```
[28]: # rounding column with scores

only_shows['imdb_score'] = only_shows['imdb_score'].round()
#checking the outcome with tail()
only_shows.head()
```

/tmp/ipykernel_26/2130509769.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
only_shows['imdb_score'] = only_shows['imdb_score'].round()
```

```
[28]:
```

	name	character	role	title	type	\
1661	Jeff Probst	Himself - Host	ACTOR	Survivor	show	
2073	Mayumi Tanaka	Monkey D. Luffy (voice)	ACTOR	One Piece	show	
2074	Kazuya Nakai	Roronoa Zoro (voice)	ACTOR	One Piece	show	
2075	Akemi Okamura	Nami (voice)	ACTOR	One Piece	show	
2076	Kappei Yamaguchi	Usopp (voice)	ACTOR	One Piece	show	

	release_year	genres	\
1661	2000	['reality']	
2073	1999	['animation', 'action', 'comedy', 'drama', 'fa...	
2074	1999	['animation', 'action', 'comedy', 'drama', 'fa...	
2075	1999	['animation', 'action', 'comedy', 'drama', 'fa...	
2076	1999	['animation', 'action', 'comedy', 'drama', 'fa...	

	imdb_score	imdb_voted
1661	7.0	24687.0
2073	9.0	117129.0
2074	9.0	117129.0
2075	9.0	117129.0
2076	9.0	117129.0

<div class="alert alert-success"; style="border-left: 7px solid green"> Reviewer's comment, v. 1

All the transformations were performed absolutely correctly

It is now time to identify outliers based on the number of votes.

```
[29]: # Use groupby() for scores and count all unique values in each group, print the result
score_count = only_shows.groupby('imdb_score')['imdb_voted'].count()
score_count
```

```
[29]: imdb_score
2.0    24
3.0    27
4.0   180
5.0   592
6.0  2494
7.0  4706
8.0  4842
9.0   557
10.0    8
Name: imdb_voted, dtype: int64
```

<div class="alert alert-success"; style="border-left: 7px solid green"> Reviewer's comment, v. 1

The dataframe was filtered and grouped flawlessly

Based on the aggregation performed, it is evident that scores 2 (24 voted shows), 3 (27 voted shows), and 10 (only 8 voted shows) are outliers. There isn't enough data for these scores for the average number of votes to be meaningful.

To obtain the mean numbers of votes for the selected scores (we identified a range of 4-9 as acceptable), use conditional filtering and grouping.

```
[30]: # filter dataframe using two conditions (scores to be in the range 4-9)
filtered_dataframe = only_shows[(only_shows['imdb_score'] >= 4) &
    ↳(only_shows['imdb_score'] <= 9)]

# group scores and corresponding average number of votes, reset index and print
    ↳the result

grouped_scores = filtered_dataframe.groupby('imdb_score')['imdb_voted'].mean().
    ↳reset_index()
grouped_scores
```

```
[30]:   imdb_score   imdb_voted
0         4.0   5277.583333
1         5.0   3143.942568
2         6.0   3481.717322
3         7.0   8727.068211
4         8.0   30299.460967
5         9.0  126904.109515
```

Now for the final step! Round the column with the averages, rename both columns, and print the dataframe in descending order.

```
[31]: # round column with averages
grouped_scores['imdb_voted'] = grouped_scores['imdb_voted'].round()

# rename columns
rename_columns = grouped_scores.rename(columns={'imdb_score':'scores',
    ↳'imdb_voted':'votes'})

# print dataframe in descending order
descending_order = rename_columns.sort_values(by= 'votes', ascending = False)
descending_order
```

```
[31]:   scores   votes
5     9.0 126904.0
4     8.0  30299.0
3     7.0   8727.0
0     4.0   5278.0
```

2	6.0	3482.0
1	5.0	3144.0

The assumption matches the analysis: the shows with the top 3 scores have the most amounts of votes.

<div class="alert alert-success"; style="border-left: 7px solid green"> Reviewer's comment, v. 1
Great! Also correct rounding and grouping in this section

0.6 Conclusion

The research done confirms that highly-rated shows released during the "Golden Age" of television also have the most votes. While shows with score 4 have more votes than ones with scores 5 and 6, the top three (scores 7-9) have the largest number. The data studied represents around 94% of the original set, so we can be confident in our findings.

<div class="alert alert-success"; style="border-left: 7px solid green"> Reviewer's comment, v. 1

Overall conclusion is an important part, where we should include the summary of the outcomes of the project.